

Anonymous Long-Video QA Artifact

Technical document for reproducibility and anonymous submission

1. Problem and artifact scope

The task is open-ended question answering over short films. Each example contains a movie identifier, a natural-language question, and an answer that must be submitted without exposing evaluation labels. The central engineering problem is long-range evidence integration: a question may depend on a name introduced early in the film, a relation that changes later, or the final state of an event chain.

This artifact releases inference and submission code only. Raw videos, subtitles, model weights, generated predictions, credentials, and evaluation records are external inputs. The boundary supports anonymous review and prevents accidental disclosure of test information.

2. Movie-level answer-bank construction

The main runner groups questions by `video_id`. For each movie it builds one request containing a timestamped transcript and chronological visual observations. All questions for that movie are placed in the same request. The model returns exactly one JSON object whose keys are the original question IDs and whose values are concise English answers.

The computation is:

$$(E_m, Q_m) \longrightarrow \{a_{m,1}, a_{m,2}, \dots, a_{m,n_m}\},$$

where E_m is shared evidence for movie m , Q_m is its complete question set, and the answer bank is emitted in one autoregressive context. The prompt asks for consistent names, aliases, relationships, chronology, causes, repeated objects, and ending state. It forbids facts unsupported by the transcript or frames.

The evidence contract is explicit. Transcript cues are serialized with start timestamps. Frames are sampled chronologically, resized to a fixed maximum side, encoded as JPEG, and placed after the text prompt with timestamp markers. Long transcripts use a head-and-tail bound so setup and ending evidence remain visible.

3. Reproducibility and failure handling

The runner accepts an OpenAI-compatible Responses endpoint and reads its API key from a process environment variable. The key is never written to output files or logs. Movie requests can be retried for transient provider failures; answers are persisted only after the complete JSON object has been parsed and validated against every question ID in that movie. Missing credentials, malformed JSON, missing answers, and incomplete movie banks are errors.

The all-or-nothing movie boundary prevents partial requests from silently mixing answers produced under different evidence or decoding conditions. A separate conversion script maps a validated answer bank into the competition schema. The audit script checks exact ID coverage, duplicate rows, empty predictions, and the required header.

4. Competition-facing compliance

The release follows the operational conventions used by the SLoMO/EvalAI workflow:

- use the live challenge phase and current metadata fields;
- submit exactly `question_id,prediction`;
- include every required ID exactly once with a non-empty prediction;
- keep public and private annotations separate;
- do not use answer keys, options, private labels, or leaderboard output during inference.

The repository starts from a fresh anonymous Git history. Identity-bearing material is outside the tree: account names, team names, local absolute paths, API tokens, private data, generated answer banks, and submission IDs. The live competition page and EvalAI portal remain authoritative for deadlines, availability, metadata, and upload policy:

<https://slomo-workshop.github.io/eccv2026/>
<https://eval.ai/>

5. Known limitations

The method depends on transcript quality, frame coverage, model endpoint availability, and context budget. This artifact does not claim that every endpoint or model identifier remains available. It does not package competition data or an official score. Reproduction requires authorized data access and an independent local audit of the generated CSV.