

Joint Movie Context for Long-Video QA

Anonymous ECCV 2026 Submission

Paper ID #00000

Abstract. This artifact describes a movie-level answer-bank pipeline for open-ended long-video question answering. The system groups questions from one film, exposes shared timestamped and visual evidence, and emits one validated JSON answer bank. Videos, labels, generated answers, credentials, and submission records remain external inputs.

Keywords: long-video understanding · movie question answering · reproducibility · anonymous artifact

1 Task and Artifact Boundary

The task is open-ended question answering over short films. Questions may require an early name, a changing relation, or the final state of an event chain. The movie is the context unit. The release contains source code, tests, and documentation only; data, videos, weights, answers, credentials, account identifiers, leaderboard records, and absolute paths remain external inputs [2, 3].

2 Movie-Level Inference Contract

For each movie m , the runner constructs one request:

$$(E_m, Q_m) \longrightarrow \{a_{m,1}, a_{m,2}, \dots, a_{m,n_m}\}, \quad (1)$$

where E_m is shared evidence and Q_m is the complete question set. The model returns one JSON object keyed by the original question IDs.

The prompt requires consistent names, aliases, relationships, event order, causes, repeated objects, and ending state. Frames are chronological JPEGs with timestamp markers; long transcripts use a head-and-tail bound.




Fig. 1: Historical official score chain. The discontinuity appears at the movie-batched multimodal source; later parent-preserving changes are smaller. Values are correct predictions out of 538.

3 Validation and Failure Handling

The runner reads credentials only from a process environment variable. A movie result is persisted only after the response parses as JSON and contains every requested question ID with a non-empty answer. Partial movie responses are never silently merged with another protocol.

The conversion and audit scripts enforce this upload contract [1]:

- exact header: `question_id,prediction`;
- every required question ID appears exactly once;
- all predictions are non-empty;
- public and private annotations remain separate;
- inference does not read answer keys, options, private labels, or leaderboard output.

4 Evidence-Bound Diagnostics

The plots are generated from `figures/figure_data.json` by a deterministic script. Figure 2 reports a frozen six-movie, 60-question factor audit: joint answering reaches 27/60, independent calls 18/60, no frames 24/60, Target16 29/60, and Dense48 30/60. The diagnostics support joint context and bound claims about simple frame changes.




Fig. 2: Frozen factor and visual-ceiling audit on six movies. Blue is the direct reference; orange bars are controlled variants. Scores are local diagnostics, not official leaderboard results.

5 Reproducibility and Anonymous Submission

The repository starts from a fresh Git history with anonymous author metadata. The ECCV review template replaces author and affiliation fields with a paper-ID placeholder and enables line numbering. Replace the placeholder only with the assigned ID before submission; do not add author names or identifying links.

The live SLoMO page remains authoritative for deadlines, track rules, and artifact policy [3]. Reproduction requires authorized data and model access followed by the repository schema audit.

References

1. EvalAI: Evalai evaluation platform. <https://eval.ai/> (2026), submission interface and evaluation service [2](#)
2. SF20K Dataset: Sf20k movie question answering dataset. <https://huggingface.co/datasets/rghermi/sf20k> (2026), dataset landing page [1](#)
3. SLoMO Workshop: Slomo workshop and challenge page. <https://slomo-workshop.github.io/eccv2026/> (2026), accessed 2026-09-05 [1](#), [2](#)